

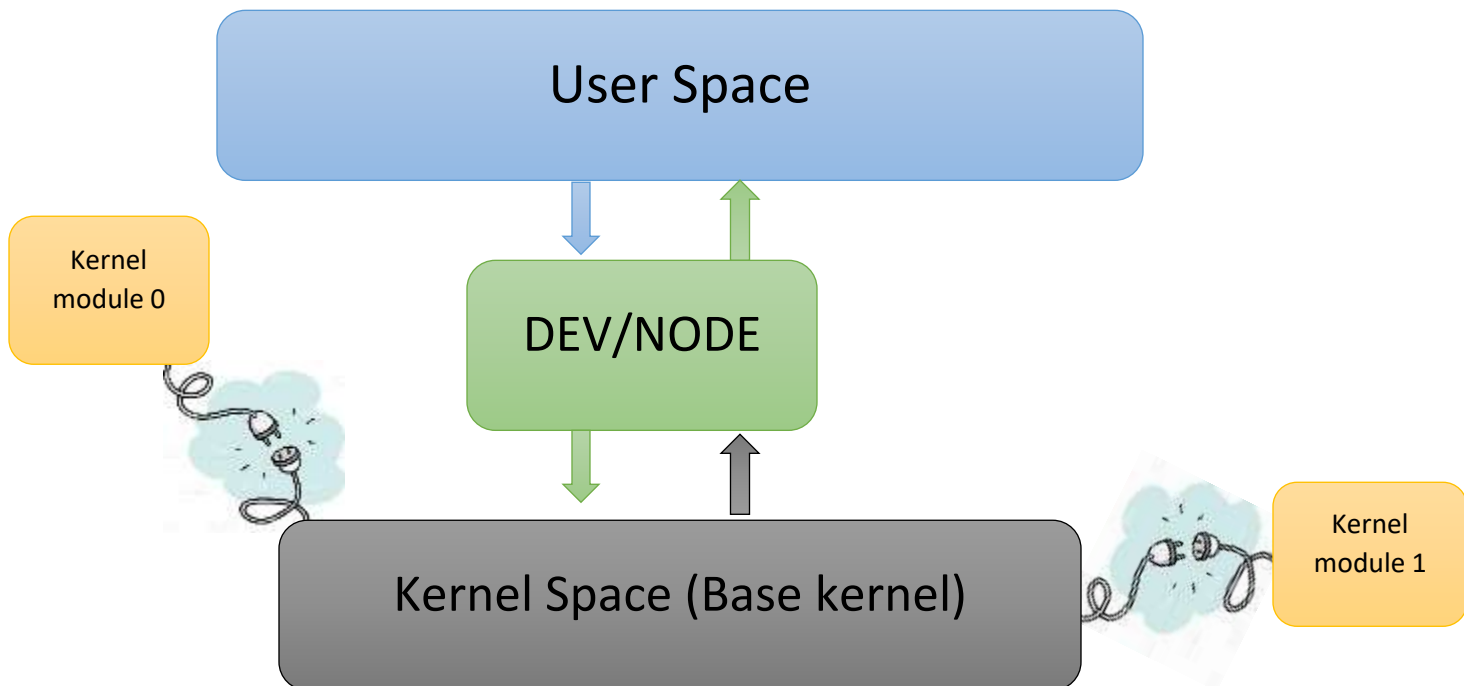
This chapter will discuss on how to

- Make a kernel module
- Implement Timer events
- And implement Interrupts in Kernel space.

And furthermore, Development of a user application (In User Space) to cater with the kernel through IOCTL (Input/output Control) commands.

For easiness of implementation let's target on Ubuntu which is a Linux kernel bases Operating system.

### Overview of Kernel Space and the User space



### Let's look at the Program Structure of a kernel Module

[Includes which would be needed]

```
#include <linux/module.h>
#include <linux/fs.h>
#include <asm/uaccess.h>
```

### [Definitions and Declarations]

```
#define KBD_IRQ      1    /* IRQ number for keyboard (i8042) */  
#define KBD_DATA_REG 0x60  
struct cdev my_cdev;  
static int  majorNumber;  
  
static int  my_open( struct inode *, struct file *);  
static ssize_t my_read( struct file *, char *, size_t, loff_t *);etc
```

### [Initializing routine]

All the initialization should be here  
This is going to be invoked during the plugging of the module to the kernel

### [Other routines]

Including subroutines, Interrupt Handlers, Callback routines etc.

### [Finalizing routine]

All used memories, interrupt requests should release and should ready to be unplugged from the kernel

Finally the Initializing routine and finalizing routines are called.

```
module_init(<Initializing routine>); //these are kernel macros  
module_init(<finalizing routine>);
```

## Procedure to make a device driver

First of all you have to have a Major number and a Minor for your device driver in order to identify the device, this can be dynamically allocated or even can be allocated manually.

Then you have to have a device class and a Device name.

Thereafter these class and the device should be registered.

```
// Try to dynamically allocate a major number for the device -- more difficult but worth it
majorNumber = register_chrdev(0, DEVICE_NAME, &my_fops); // 0 wil auto allocate major num
if (majorNumber<0) {
    printk(KERN_ALERT "char driver failed to register a major number\n");
    return majorNumber;
}
printk(KERN_INFO "char driver: Major number %d\n", majorNumber);

// Register the device class
CharClass = class_create(THIS_MODULE, CLASS_NAME);
if (IS_ERR(CharClass)){           // Check for error and clean up if there is
    unregister_chrdev(majorNumber, DEVICE_NAME);
    printk(KERN_ALERT "Failed to register device class\n");
    return PTR_ERR(CharClass);    // Correct way to return an error on a pointer
}

// Register the device driver
CharDevice = device_create(CharClass, NULL, MKDEV(majorNumber, 0), NULL, DEVICE_NAME);
if (IS_ERR(CharDevice)){         // Clean up if there is an error
    class_destroy(CharClass);    // Repeated code but the alternative is goto statements
    unregister_chrdev(majorNumber, DEVICE_NAME);
    printk(KERN_ALERT "Failed to create the device\n");
    return PTR_ERR(CharDevice);
}
```

Since we are focusing on a char driver, there should be a way to communicate with the user space, for that we use the file operation structure in kernel.

```
struct file_operations my_fops = {
    read  :   my_read,
    write :   my_write,
    open  :   my_open,
    release : my_close,
    unlocked_ioctl : my_ioctl,
    owner  :   THIS_MODULE
};
```

In here user space have the access to the read, write ... functions. The arguments are defined in the kernel space functions which are related as above.

```
static ssize_t my_read(struct file *filp, char *buff, size_t len, loff_t *off)
{
    short count;
    printk("\n*****Some body is Reading me*****\n");
```

```

        count = copy_to_user(buff, msg, len);
        return 0;
    }

```

\*Here the buff means a char pointer to the user space buffer and the msg means the kernel space buffer

## Implementing an Interrupt

In the process of implementing an interrupt for the kernel module 4 things should be aware of,

- Selecting the interrupt request number (IRQ number)
- Requesting for an interrupt
- Implementing the interrupt handler also named as Interrupt service routine (ISR)
- Finally make sure to release/free the IRQ

```

    • #define KBD_IRQ      1    //IRQ number for the keyboard
    • request_irq(KBD_IRQ, kbd2_isr, IRQF_SHARED, "kbd2", (void *)kbd2_isr);
    • static irqreturn_t kbd2_isr(int irq, void *dev_id){
        //your ISR code goes here
        return IRQ_HANDLED;
    }
    • free_irq(KBD_IRQ, (void *)kbd2_isr);

```

## Implementing an Timer callback event

In the process of implementing a Timer callback event for the kernel module few things should be aware of,

### Declaring

```

static struct timer_list my_timer;
void my_timer_callback( unsigned long data );

```

### Initializing

```

setup_timer(&my_timer, my_timer_callback, 0);/* setup your timer to call my_timer_callback */
mod_timer(&my_timer, jiffies + msecs_to_jiffies(200));/* setup timer interval to 200 msecs */

```

### Timer callback routine

```

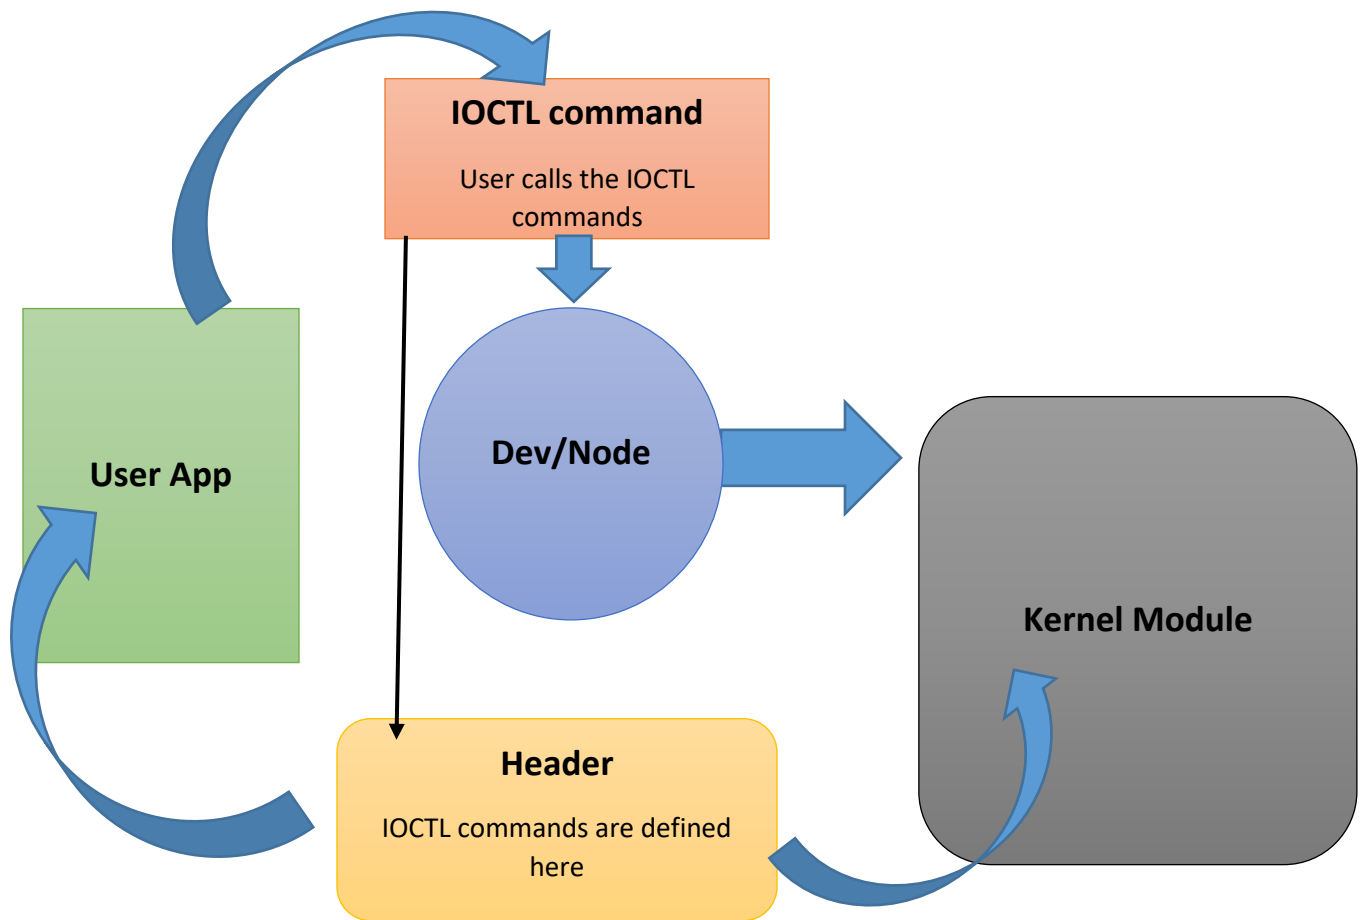
void my_timer_callback( unsigned long data )
{
    mod_timer(&my_timer, jiffies + msecs_to_jiffies(2000));/* setup timer interval to 2000 msecs */
    printk(KERN_ALERT "2 seconds Gone \n");
}

```

## Deleting the Timer

`del_timer(&my_timer);` /\* remove kernel timer when unloading module\*/

### How IOCTL works



## User program

```
switch(input){  
    case 0:  
        ret_val = ioctl(fd, IOCTL_CMD0, 0);  
        if (ret_val < 0){  
            printf("Failed IOCTL_CMD0. %d\n",  
ret_val);  
        }  
        break;  
}
```

## Kernel Program

```
static long my_ioctl(struct file *file, unsigned int  
command, unsigned long value)  
{  
    switch(command){  
        case IOCTL_CMD0:  
            printk(KERN_INFO "I am CMD0\n");  
            break;  
    }  
    return 0;  
}
```

## Header

```
#define IOCTL_CMD0 _IOW(Magic_Num, 0, int)
```

**IOCTL makes it user friendly for the user app to communicate with the kernel module**

Example code to capture keyboard press, evoke a timer event and User application to communicate with the kernel module.

## User Application

```
#include <stdio.h>  
#include <string.h>  
#include <fcntl.h>  
#include <linux/ioctl.h>  
#include <sys/ioctl.h>  
#include <unistd.h>  
  
#include "User.h"  
  
int main(){  
  
    int input = 0;  
    char buff;  
    int fd;  
    int ret_val=0;  
  
    fd = open("/dev/TEST_DRIVER", O_RDWR);  
  
    if(fd == 0){  
        printf("Error opening file");  
    }
```

```

}else{

    printf("Enter 1/0 : ");
    scanf("%d", &input);
    //write(fd,&input,1);
    switch(input){
        case 0:
            ret_val =ioctl(fd,IOCTL_CMD0, 0);
            if (ret_val < 0){
                printf("Failed IOCTL_CMD0: %d\n", ret_val);

            }
            break;
        case 1:
            ret_val = ioctl(fd, IOCTL_CMD1, 0);
            if (ret_val < 0){
                printf("Failed IOCTL_CMD1: %d\n", ret_val);

            }
            break;
        case 2:
            ret_val = ioctl(fd, IOCTL_CMD2,&buff);
            if (ret_val < 0){
                printf("Failed IOCTL_CMD2: %d\n", ret_val);

            }
            printf("Buffer : %d\n", buff);
            break;
        default:
            printf("Error input");

    }

    close(fd);

}

return 0;
}

```

## **Header**

```

#ifndef TEST_H

#define TEST_H

#define MY_MAJOR 200
#define MY_MINOR 0
#define MY_DEV_COUNT 1
#define DEVICE_NAME "TEST_DRIVER"    ///< The device will appear at /dev
#define CLASS_NAME "TDRER"          ///< The device class -- this is a character device driver

#include <linux/ioctl.h>

```

```

#define Magic_Num 131

#define IOCTL_CMD0_IOW(Magic_Num, 0, int)
#define IOCTL_CMD1_IOW(Magic_Num, 1, int)
#define IOCTL_CMD2_IOR(Magic_Num, 2, char*)

#define DEVICE_FILE_NAME "TEST_DRIVER"

#endif

```

## **Kernel Module**

```

/*
 * This is a keyboard tracking kernal module;
 * once this module is pulgged into the kernal
 * all the track recods are reserved in the kernal.
 *
 * Devoloped by: Malinda Sulochana Silva
 * Organization: Zone24x7
 *
 * Copyright. All Rights Reserved.
 * */

#include <linux/module.h>
#include <linux/string.h>
#include <linux/fs.h>
#include <asm/uaccess.h>
#include <linux/init.h>
#include <linux/cdev.h>
#include <linux/device.h>
#include <linux/kernel.h>
#include <linux/errno.h>
#include <linux/io.h>
#include <linux/sched.h>
#include <linux/interrupt.h>
#include <linux/list.h>
#include <linux/irq.h>
#include <linux/slab.h>
#include <linux/gpio.h>
#include <linux/time.h>
#include <linux/timer.h>
#include <linux/delay.h>
#include <linux/ioctl.h>
#include "User.h"

// _____ Definitions

#define KBD_IRQ      1      /* IRQ number for keyboard (i8042) */
#define KBD_DATA_REG 0x60   /* I/O port for keyboard data */
#define KBD_SCANCODE_MASK 0x7f
#define KBD_STATUS_MASK 0x80

MODULE_LICENSE("GPL");
MODULE_AUTHOR("Malinda Sulochana Silva");

```



```
MODULE_DESCRIPTION("A Simple GPIO Device Driver module for Ubuntu");
```

```
//_____Declarations
```

```
unsigned long j, stamp_1, stamp_half, stamp_n;
```

```
int malinda=123;
```

```
int n=2;
```

```
static char *msg=NULL;
```

```
struct cdev my_cdev;
```

```
static int majorNumber;
```

```
static int my_open( struct inode *, struct file *);
```

```
static ssize_t my_read( struct file *, char *, size_t, loff_t *);
```

```
static ssize_t my_write(struct file *, const char *, size_t, loff_t *);
```

```
static int my_close(struct inode *, struct file *);
```

```
static long my_ioctl(struct file *, unsigned int, unsigned long);
```

```
static irqreturn_t kbd2_isr(int irq, void *dev_id);
```

```
static struct timer_list my_timer;
```

```
void my_timer_callback( unsigned long data );
```

```
static struct class* CharClass = NULL; ///< The device-driver class struct pointer
```

```
static struct device* CharDevice = NULL; ///< The device-driver device struct pointer
```

```
/*+++++++_MAIN_+++++++*/
```

```
/*|||||*/
```

```
/*_____File Operations_____*/
```

```
struct file_operations my_fops = {
```

```
    read : my_read,
```

```
    write : my_write,
```

```
    open : my_open,
```

```
    release : my_close,
```

```
    unlocked_ioctl : my_ioctl,
```

```
    owner : THIS_MODULE
```

```
};
```

```
/*=====*/
```

```
/*-----*/
```

```
/*_____File Operation Routines_____*/
```

```
static int my_open(struct inode *inode, struct file *fil)
```

```
{
```

```
    printk("\n*****Some body is opening me*****\n");
```

```
    return 0;
```

```
}
```

```
static ssize_t my_read(struct file *filp, char *buff, size_t len, loff_t *off)
```

```
{
```

```
    short count;
```

```
    printk("\n*****Some body is Reading me*****\n");
```

```
    count = copy_to_user(buff, msg, len);
```

```
    return 0;
```

```

}
static ssize_t my_write(struct file *filp, const char *buff, size_t len, loff_t *off)
{
    short count;
    memset(msg, 0, 32);
    printk("\n*****Some body is writting to me*****\n");
    // -- copy the string from the user space program which open and write this device
    count = copy_from_user( msg, buff, len );
    printk("%s\n",msg);
    return count;
}
static int my_close(struct inode *inode, struct file *fil)
{
    printk("\n*****Some body is Closing me*****\n");
    return 0;
}
/*=====*/

/*-----*/
/* _____IOCTL Routine_____ */
static long my_ioctl(struct file *file, unsigned int command, unsigned long value)
{
    switch(command){

        case IOCTL_CMD0:
            printk(KERN_INFO "I am CMD0\n");
            break;

        case IOCTL_CMD1:
            printk(KERN_INFO "I am CMD1\n");
            break;

        case IOCTL_CMD2:
            printk(KERN_INFO "I am CMD2\n");
            break;

        default:
            printk(KERN_INFO "I am default");
            break;
    }

    return 0;
}
/*=====*/

/*-----*/
/* _____Timer_____ */
void my_timer_callback( unsigned long data )
{

```

```

        mod_timer(&my_timer, jiffies + msecs_to_jiffies(2000));/* setup timer interval to 2000 msecs */
        printk(KERN_ALERT "2 seconds Gone \n");

    }

/*=====*/

/*-----*/
/* _____Interrupt Handler_____ */

static irqreturn_t kbd2_isr(int irq, void *dev_id)
{
    char scancode;
    char key1=0;
    scancode = inb(KBD_DATA_REG);
    if((int)scancode==30){key1='A';}
    else if((int)scancode==48){key1='B';}
    else if((int)scancode==46){key1='C';}
    else if((int)scancode==32){key1='D';}
    else if((int)scancode==18){key1='E';}
    else if((int)scancode==33){key1='F';}
    else if((int)scancode==34){key1='G';}
    else if((int)scancode==35){key1='H';}
    else if((int)scancode==23){key1='I';}
    else if((int)scancode==36){key1='J';}
    else if((int)scancode==37){key1='K';}
    else if((int)scancode==38){key1='L';}
    else if((int)scancode==50){key1='M';}
    else if((int)scancode==49){key1='N';}
    else if((int)scancode==24){key1='O';}
    else if((int)scancode==25){key1='P';}
    else if((int)scancode==16){key1='Q';}
    else if((int)scancode==19){key1='R';}
    else if((int)scancode==31){key1='S';}
    else if((int)scancode==20){key1='T';}
    else if((int)scancode==22){key1='U';}
    else if((int)scancode==47){key1='V';}
    else if((int)scancode==17){key1='W';}
    else if((int)scancode==45){key1='X';}
    else if((int)scancode==21){key1='Y';}
    else if((int)scancode==44){key1='z';}
    else if((int)scancode==14){key1='<';}
    else if((int)scancode==28){key1='|';}
    else if((int)scancode==57){key1='_';}
    else if((int)scancode==57){key1=':';}
    /* NOTE: i/o ops take a lot of time thus must be avoided in HW ISRs */
    //just like printk included with KERN_ALERT
    //pr_info("Scan Code %c %s\n", key1 & KBD_SCANCODE_MASK,scancode & KBD_STATUS_MASK ?
                                                    "Released" : "Pressed");

    pr_info("Scan Code %c \n", key1 & KBD_SCANCODE_MASK);

    return IRQ_HANDLED;
}

```

```

}
/*=====*/

/*-----*/
/* _____Driver Initialization_____ */
int Start(void)
{
    // Try to dynamically allocate a major number for the device -- more difficult but worth it
    majorNumber = register_chrdev(0, DEVICE_NAME, &my_fops); //auto allocate major num
    if (majorNumber<0) {
        printk(KERN_ALERT "char driver failed to register a major number\n");
        return majorNumber;
    }
    printk(KERN_INFO "char driver: Major number %d\n", majorNumber);

    // Register the device class
    CharClass = class_create(THIS_MODULE, CLASS_NAME);
    if (IS_ERR(CharClass)){          // Check for error and clean up if there is
        unregister_chrdev(majorNumber, DEVICE_NAME);
        printk(KERN_ALERT "Failed to register device class\n");
        return PTR_ERR(CharClass);    // Correct way to return an error on a pointer
    }

    // Register the device driver
    CharDevice = device_create(CharClass, NULL, MKDEV(majorNumber, 0), NULL, DEVICE_NAME);
    if (IS_ERR(CharDevice)){          // Clean up if there is an error
        class_destroy(CharClass);      // Repeated code but the alternative is goto statements
        unregister_chrdev(majorNumber, DEVICE_NAME);
        printk(KERN_ALERT "Failed to create the device\n");
        return PTR_ERR(CharDevice);
    }
    printk(KERN_INFO "char driver: device class created correctly\n"); // Made it! device was initialized

    /* _____Registering an Interrupt Handler_____ */
    request_irq(KBD_IRQ, kbd2_isr, IRQF_SHARED, "kbd2", (void *)kbd2_isr);

    /* _____Setting an Timer Callback event_____ */
    setup_timer(&my_timer, my_timer_callback, 0);/* setup your timer to call my_timer_callback */
    mod_timer(&my_timer, jiffies + msecs_to_jiffies(200));/* setup timer interval to 200 msecs */

    return 0;
}
/*=====*/

/*-----*/
/* _____Driver Disposing_____ */

void Dispose(void)
{

```

```

device_destroy(CharClass, MKDEV(majorNumber, 0)); // remove the device
class_unregister(CharClass); // unregister the device class
class_destroy(CharClass); // remove the device class
unregister_chrdev(majorNumber, DEVICE_NAME); // unregister the major number
printk(KERN_INFO "TEST_DRIVER char driver: Goodbye \n");
free_irq(KBD_IRQ, (void *)kbd2_isr);
del_timer(&my_timer);/* remove kernel timer when unloading module*/
}
/*=====*/

```

```

module_init( Start );
module_exit( Dispose );

```

## OUTPUTS:

The image shows a terminal window titled 'malinda@ubuntu: ~/Desktop/KernalModules'. The output of the kernel module is displayed in green text on a black background. Several lines are highlighted in red, indicating specific events. Two blue callout boxes with yellow arrows point to these events: 'Timer callback event' points to the '2 seconds Gone' messages, and 'Keyboard press interrupt' points to the 'Scan Code' messages.

```

malinda@ubuntu: ~/Desktop/KernalModules
[ 1087.776381] char driver: device class created correctly
[ 1087.899678] Scan Code
[ 1087.972841] 2 seconds Gone
[ 1089.115642] Scan Code S
[ 1089.259646] Scan Code
[ 1089.974993] 2 seconds Gone
[ 1090.747637] Scan Code <
[ 1090.867618] Scan Code
[ 1091.976702] 2 seconds Gone
[ 1093.978009] 2 seconds Gone
[ 1095.980347] 2 seconds Gone
[ 1096.067621] Scan Code D
[ 1096.179624] Scan Code
[ 1096.595632] Scan Code M
[ 1096.675480] Scan Code
[ 1096.787565] Scan Code E
[ 1096.851520] Scan Code
[ 1097.043599] Scan Code S
[ 1097.139624] Scan Code
[ 1097.771641] Scan Code G
[ 1097.851574] Scan Code
[ 1097.982523] 2 seconds Gone
[ 1098.235519] Scan Code |
malinda@ubuntu:~/Desktop/KernalModules$

```

Annotations:

- Timer callback event (points to '2 seconds Gone' lines)
- Keyboard press interrupt (points to 'Scan Code' lines)

Read, write, and execute permissions

```
malinda@ubuntu:~/Desktop/KernalModules$ sudo insmod module1.ko
malinda@ubuntu:~/Desktop/KernalModules$ cd User/
malinda@ubuntu:~/Desktop/KernalModules/User$ gcc User.c
malinda@ubuntu:~/Desktop/KernalModules/User$ sudo chmod 777 /dev/TEST_DRIVER
malinda@ubuntu:~/Desktop/KernalModules/User$ ls -la /dev/TEST_DRIVER
crwxrwxrwx 1 root root 250, 0 Nov 17 01:09 /dev/TEST_DRIVER
malinda@ubuntu:~/Desktop/KernalModules/User$ ./a.out
Enter 1/0 : 1
malinda@ubuntu:~/Desktop/KernalModules/User$ ./a.out
Enter 1/0 : 0
malinda@ubuntu:~/Desktop/KernalModules/User$
```

Open routine evoked from the user space

IOCTL command 1

IOCTL command 0

```
malinda@ubuntu:~/Desktop/KernalModules/User
[ 1447.883666] Scan Code
[ 1447.915587] Scan Code T
[ 1448.043641] Scan Code
[ 1448.498838] 2 seconds Gone
[ 1448.619557] Scan Code |
[ 1448.622075]
*****Some body is opening me*****
[ 1448.747715] Scan Code
[ 1449.915652] Scan Code
[ 1450.027661] Scan Code
[ 1450.503463] 2 seconds Gone
[ 1450.715542] Scan Code |
[ 1450.717574] I am CMD1
[ 1450.717579]
*****Some body is Closing me*****
[ 1450.859571] Scan Code
[ 1452.507075] 2 seconds Gone
[ 1452.715601] Scan Code
[ 1452.715790] Scan Code
[ 1452.716173] Scan Code
[ 1452.716411] Scan Code
[ 1452.827592] Scan Code
[ 1452.827793] Scan Code
[ 1452.828214] Scan Code
[ 1452.828665] Scan Code
[ 1453.523533] Scan Code |
[ 1453.526194]
*****Some body is opening me*****
[ 1453.635637] Scan Code
[ 1454.511099] 2 seconds Gone
[ 1456.315612] Scan Code
[ 1456.435626] Scan Code
[ 1456.515415] 2 seconds Gone
[ 1457.115523] Scan Code |
[ 1457.117627] I am CMD0
[ 1457.117633]
*****Some body is Closing me*****
[ 1457.243637] Scan Code
[ 1458.519737] 2 seconds Gone
[ 1460.523588] 2 seconds Gone
[ 1462.526960] 2 seconds Gone
```



## Compiling the kernel Module and Makefile

In order to compile the kernel module doing “`gcc module.c`” simply doesn’t work. Because to build the kernel module it needs special header files, which we called as **kernel headers**. And especially these headers should be compatible with the relevant base kernel which you are going to plug your modules into.

What you have to do is, simply make a “**Makefile**” as follows

```
obj-m += module1.o

KDIR := /usr/src/linux-headers-3.19.0-15-generic
PWD:= $(shell pwd)
default:
    $(MAKE) -C $(KDIR) M=$(PWD) modules
clean:
    $(MAKE) -C $(KDIR) M=$(PWD) clean
```

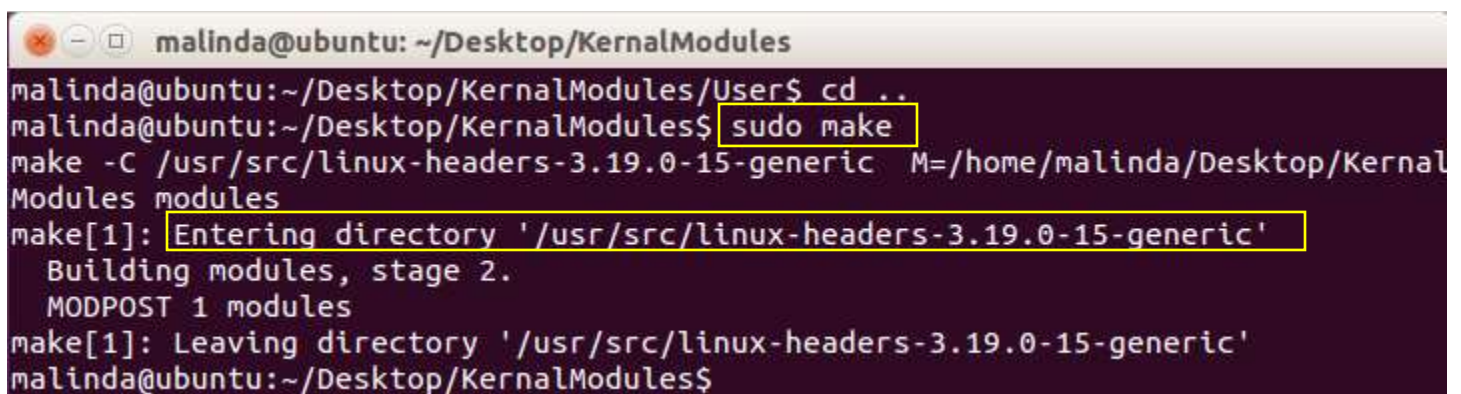
Here KDIR means the directory where your kernel headers are located

PWD is where your kernel module exist (i.e module.c)

Thereafter what you have to do is type **sudo make** in the terminal

What it does is it will find a file named as (**Makefile**) inside the current directory and compile it

So, if the compilation is success you will get new bunch of files one having a name “**<module.ko>**”. And that is your kernel object file. Now you just have to plug it into the kernel and test for it



```
malinda@ubuntu: ~/Desktop/KernalModules
malinda@ubuntu:~/Desktop/KernalModules/Users$ cd ..
malinda@ubuntu:~/Desktop/KernalModules$ sudo make
make -C /usr/src/linux-headers-3.19.0-15-generic M=/home/malinda/Desktop/KernalModules modules
make[1]: Entering directory '/usr/src/linux-headers-3.19.0-15-generic'
Building modules, stage 2.
MODPOST 1 modules
make[1]: Leaving directory '/usr/src/linux-headers-3.19.0-15-generic'
malinda@ubuntu:~/Desktop/KernalModules$
```

When compiling the user program you just have type `"gcc User.c"` and it will create a file `"a.out"`  
Thereafter to run the file as usual  
Just type `"./a.out"`

### Things you should get into know about (Terminal tips)

- The node is create/created in the , `/dev/< node_name >`
- Make node manually `sudo mknod /dev/<node_name> c majorNo MinorNo`
- Taking read, write and execute permission `sudo chmod 777 /dev/<node_name>`
- Look into the kernel space messages `dmesg`
- Go to previous Directory `cd ..`
- List files `ls`
- Inserting a module `sudo insmod <module.ko>`
- Remove a module `sudo rmmod <module>`
- Check available modules `lsmod`